

TENNIS MATRIX — DISAGREEMENT / TRAP AUDIT

UPDATED WITH NEW VERIFIED WIN/LOSS LEARNING + CURRENT CALIBRATION RULES

NEW POST-RESULT LEARNING / 120-MATCH CALIBRATION ENFORCEMENT

These rules are additive and override any earlier rule that would allow a candidate to bypass them.

1. UPDATED CALIBRATION POSITION

The current verified calibration/backtest must run AFTER the Verification Audit, Disagreement/Trap Audit, Dangerous-Underdog screen, stress tests, and removal case, but BEFORE the final color/output. Recompute from the newest verified sample every time it changes. Never reuse stale bucket rates.

2. NO-EDGE / LOW-PROBABILITY FLOOR

- Headline or independently reconstructed favorite probability $\leq 55\%$: automatic PASS/RED; never Green.
- 56-64%: cannot be Green unless the favorite defeats every major underdog pathway with at least two independent high-reliability evidence families and no unresolved critical contradiction.
- A Red/PASS classification means "selection is not reliable enough"; it does NOT automatically mean flip to the opponent.

3. INDEPENDENT FLIP STANDARD

Never FLIP solely because of one contradictory family, recent form alone, opponent-adjustment alone, H2H alone, ranking alone, or generic external narrative.

A FLIP requires:

- (a) at least 2 genuinely independent high-reliability underdog evidence families;
- (b) at least 1 independent external confirmation;
- (c) the underdog survives component-removal and conservative stress testing;
- (d) the favorite fails the strongest evidence-based defense.

If these are not met, DOWNGRADE/PASS rather than FLIP.

4. DANGEROUS UNDERDOG CEILING OVERRIDE

For every underdog, explicitly search for a recent ceiling event:

- win over Top 10/20/50 or clearly elite/high-level opponent;
- repeated competitive performances against much stronger opposition;
- rapid ranking/Elo acceleration;
- qualifying-to-main-draw or level-jump surge;
- major same-surface breakthrough;
- recent hold/break/serve/return performance materially above long-run rating;
- comeback/deciding-set/tiebreak resilience against high-quality opposition.

A qualifying ceiling event triggers mandatory escalation. It cannot be dismissed merely because the favorite leads aggregate Elo, market, Monte Carlo, or model vote count.

5. PRESENT-STRENGTH / RATING-LAG TEST

Compare long-run ratings against last 30/60/90-day and same-surface performance. Explicitly classify each player:

CURRENTLY REPRESENTATIVE / POSSIBLY STALE / CLEARLY LAGGING.

If the underdog's rating is lagging rapid improvement, historical Elo/ranking receives reduced decision authority.

6. ELITE-FAVORITE EXEMPTION PROHIBITED

No 80%, 85%, 90%+ favorite is exempt from the full underdog audit. High headline probability, Strong Agreement, Elite recommendation, or high Monte Carlo cannot suppress a demonstrated dangerous-underdog signal.

7. H2H QUALITY CONTROL

H2H must be weighted for recency, surface, tour level, age/development stage, current player level, and sample size. Old/thin/non-comparable H2H cannot veto stronger present-strength evidence.

8. CORRELATED AGREEMENT COMPRESSION

Surface Elo, General, Specialist, Monte Carlo, headline probability, and market may share inputs. Count effective independent evidence families, not raw agreeing metrics. Strong Agreement is not sufficient for Green if effective independent evidence is low.

9. GREEN / DOUBLE-GREEN EXIT STANDARD

Before Green, the favorite must:

- defeat the underdog's 5 strongest pre-match arguments;
- survive the conservative probability stress range;
- survive removal of the strongest favorite component;
- have no unresolved Dangerous-Underdog ceiling trigger;
- have sufficient CRITICAL evidence coverage;
- survive updated historical calibration.

Double Green requires all Green conditions plus exceptionally strong independent evidence and no material unresolved contradiction.

10. POST-RESULT LEARNING LEDGER

Every completed match must be appended to the verified calibration. For every Green/Double-Green loss, classify the pre-match failure mechanism using only evidence available before first serve. For every successful downgrade/removal and every false downgrade/flip, record why. Recurrent structures become future downgrade or anti-overreaction rules.

DISAGREEMENT-SPECIFIC ADDITIONS

A. FALSE-POSITIVE TRAP DEFENSE

For every proposed downgrade/flip, construct the strongest case that the apparent disagreement is noise/correlation/stale evidence. If that defense survives, do not flip.

B. DISAGREEMENT QUALITY > DISAGREEMENT COUNT

Report each contradiction as INDEPENDENT / PARTIALLY CORRELATED / DUPLICATIVE and HIGH / MEDIUM / LOW reliability. A pile of correlated disagreement cannot outweigh multiple independent favorite families.

C. FLIP VS PASS SEPARATION

The audit must output KEEP / DOWNGRADE / REMOVE-PASS / FLIP separately. Never convert "favorite unsafe" into "underdog is the pick" without satisfying the Independent Flip Standard.

D. RECENT-FORM OVERREACTION GUARD

Recent form and opponent adjustment can create CAUTION or removal, but cannot independently force a flip when Elo, serve/return, market, and other independent evidence remain materially opposite.

E. CEILING-EVENT HANDOFF

Any verified underdog ceiling event must be carried verbatim into the final trap packet and stress test. The final gate must explicitly state why the favorite still survives it, or Green is blocked.

120-MATCH EMPIRICAL DISAGREEMENT / TRAP OVERLAY

This section is additive. It updates how the existing Disagreement/Trap Audit interprets contradictions; it does not delete the original audit.

A. Governing empirical rule

Disagreement QUALITY outranks disagreement COUNT. A contradiction becomes decision-relevant only after checking independence, reliability, sample, magnitude, recency, surface relevance, and opponent adjustment. Raw vote count is not a Green or Red rule.

B. Mandatory high-risk structures

Close / coin-flip structure: current 120-match record was 14/27 (51.9%). Treat this as a major trap zone. Green requires unusually strong independent evidence and successful stress testing.

Elo separation under 50 points: current record was 32/52 (61.5%) versus 43/53 (81.1%) when the Matrix pick held an Elo gap of at least 50. A sub-50 gap is a material disagreement/fragility flag, not an automatic flip.

Monte Carlo below 55%: current record was 14/22 (63.6%). This is downgrade pressure. MC $\geq 70\%$ was 11/12, but that sample is too small to function as a standalone Green rule.

Specialist support below 65%: current record was 12/22 (54.5%). This becomes a major warning when the headline probability is materially higher.

C. Correlated-agreement penalty

General Model, Specialist Model, Surface Elo, headline probability, Monte Carlo, and Market may share inputs. Do not count them as independent simply because they all agree. Strong Model Agreement is contextual only; it cannot rescue a weak independent-family structure.

D. Metrics demoted to contextual status

Data Quality, Upset Risk, Recommendation label, raw disagreement count, fatigue alone, H2H alone, and a single Recent Form contradiction are not standalone downgrade/flip triggers. They remain reportable context.

E. Required trap decision sequence

1. Check Matchup Closeness and headline WP burden zone. 2. Measure Elo separation and Surface Elo support with reliability/sample. 3. Compare MC support and headline-vs-MC gap. 4. Test recent-form differential, but require opponent adjustment. 5. Search for dangerous-underdog ceiling events and demonstrated matchup pathways. 6. Separate independent evidence from correlated Matrix outputs. 7. Remove the favorite's strongest family and rerun. 8. Output KEEP / DOWNGRADE / REMOVE-PASS / FLIP separately.

F. Empirical combinations

Elo gap ≥ 50 + MC ≥ 55 : $40/49 = 81.6\%$. **Elo gap ≥ 50 + recent-form advantage > 10 :** $17/20 = 85.0\%$. **Close + MC < 55 :** $7/13 = 53.8\%$. **Close + Elo gap < 50 :** $11/21 = 52.4\%$. These combinations are calibration overlays, not permanent laws.

G. Green-blocking rule added

A candidate cannot receive Green solely because the Matrix headline, Recommendation, Strong Agreement, General Model, Specialist Model, and Surface Elo all point the same way. Green requires independent-family support and must survive the underdog ceiling, present-strength, matchup, and component-removal tests.

ORIGINAL AUDIT — PRESERVED IN FULL

Page 1

TENNIS MATRIX AI — DISAGREEMENT AUDIT (Reusable Template)

Purpose: For any matchup, don't ask "is the pick probably right." Actively hunt for every distinct signal in the Matrix data that disagrees with, undercuts, or contradicts the predicted winner. Rank matches by how many independent disagreement signals exist against the pick, most-contested first. A 0-disagreement match isn't "safe" in any outside sense — it just means nothing in the data pushes back on it.

Nothing in a completed audit should be invented — every line must trace to a specific field in the Matrix export or the backtest workbook. Where a category has nothing to report for a given match, leave it out of that match's list rather than padding it.

Categories to check for disagreement — every match, every time

1. Each live model component individually (Surface Elo, Serve/Return, Recent Form, Market, General Model, Specialist Model) — does it name the opponent instead of the Matrix pick?
2. H2H — check separately. If reliability is near-zero/excluded (as it typically is on thin sample), note it but don't count it as real signal.
3. Component vote breadth — does a majority of live (non-excluded) components actually back the pick, or does the pick win only because one heavily-weighted component carries it?
4. Monte Carlo —
 - Does the simulated win probability fall under 50% for the pick even though it's still the named favorite?
 - Does the simulated probability diverge sharply (≥ 15 points) from the headline win probability?
 - How far down does the simulation's downside range go (e.g., a range floor at 30% or below is worth flagging)?
5. Fatigue — is the picked player actually more fatigued (higher fatigue index) than the opponent?
6. The Matrix's own self-reported flags —
 - Upset Risk (Moderate / High / EXTREME)
 - Data Quality (flag anything under ~70%)
 - Model Agreement label (High Disagreement / Mixed — the components disagree with each other, not just with the pick)
 - Recommendation label itself (Data Incomplete, Low Confidence, etc.)
7. The historical backtest overlay — does this match's closeness tier / win-probability band / recommendation label historically underperform in the verified backtest sample? (See reference table below — recompute if the backtest workbook changes.)

Backtest reference table (recompute whenever a new/updated verified-results workbook is supplied — do not reuse old numbers on a different sample)

| Pattern | Sample | Win rate |

| ---|---|---|

| Overall | — | — |

| Closeness: Clear favorite | — | — |

| Closeness: Moderate lean | — | — |

| Closeness: Close matchup | — | — |

| WP band: 60–65% | — | — |

| WP band: 65–70% | — | — |

| WP band: 70–75% | — | — |

| WP band: 75–80% | — | — |

| WP band: 80–85% | — | — |

| WP band: 85–90% | — | — |

Page 2

| WP band: 90%+ | — | — |

| Recommendation: Elite | — | — |

| Recommendation: High Confidence | — | — |

| Recommendation: Moderate Confidence | — | — |

| Recommendation: Highest Confidence | — | — |

| Model Agreement: Strong | — | — |

| Model Agreement: Mixed | — | — |

| Model Agreement: High Disagreement | — | — |

| Model Agreement: Moderate | — | — |

| Upset Risk: Low | — | — |

| Upset Risk: Moderate | — | — |

Flag any bucket with a small sample (roughly under 10) as "suggestive only, not usable as a standalone rule."

Ranked summary table (fill in per batch of matchups)

| # Disagreements | Match | Matrix Pick | WP | Color |

|---|---|---|---|

| | | | |

Per-match entry template (copy this block for every matchup)

[Player 1] vs [Player 2] — [N] disagreement signal(s) found

Matrix pick: [Player] ([WP]% headline WP) | Current color: [Green/Yellow/Red]

- Surface Elo favors [opponent], contributing [X] pts against the pick. *(omit if it favors the pick)*
- Serve/Return favors [opponent]... *(omit if it favors the pick)*
- Recent Form favors [opponent]... *(omit if it favors the pick)*
- Market favors [opponent]... *(omit if it favors the pick)*
- General Model favors [opponent]... *(omit if it favors the pick)*
- Specialist Model favors [opponent]... *(omit if it favors the pick)*
- H2H nominally favors [opponent], but reliability is near-zero/excluded — noted, not counted as real signal. *(omit if H2H favors the pick or has real sample)*
- Vote breadth: only [X] of [Y] live components actually back [pick] — [majority/tie] favors the opponent. *(omit if unanimous)*
- Monte Carlo disagrees numerically: simulation gives [pick] only [X]% (below 50%) despite being named favorite. *(omit if MC ≥50%)*
- Headline-vs-simulation gap: headline WP is [X]% but Monte Carlo only supports [Y]% — a [gap]-point gap. *(omit if gap <15)*
- Monte Carlo downside range reaches as low as [X]% for [pick] across simulation runs. *(omit if floor >30%)*
- Fatigue favors the opponent: [pick]'s fatigue index ([X]) is worse than the opponent's ([Y]). *(omit if pick is fresher or equal)*
- Matrix's own Upset Risk flag is '[level]' — the system itself is signaling doubt. *(omit if Low)*
- Data Quality is only [X]% — too thin to fully trust the pick either way. *(omit if ≥70%)*
- Matrix's own Model Agreement label is '[label]' — the components disagree with each other. *(omit if Strong/Moderate)*
- Historical pattern: this match's closeness tier ('[tier]') hit [X]% ([n]/[N]) in the backtest. *(include the actual number from the reference table above; flag explicitly if it's a weak/underperforming tier)*
- Historical pattern: WP [X]% sits in the [band] bucket, which hit [Y]% in the backtest. *(include only if that band is notably weak)*

Page 3

- Historical pattern: recommendation label '[label]' hit [X]% in the backtest. *(include only if that label is notably weak, e.g. "High Confidence" underperforming "Moderate Confidence")*

(If no category produced a hit): No disagreement signal found in any category. Every live component, the vote count, Monte Carlo, fatigue, the Matrix's own flags, and the backtest profile all line up behind this pick.

Reading the finished audit

- A high disagreement count does not automatically mean "flip the pick" — it means the pick is surviving real pushback in the data, which is more useful to know than a clean confidence score alone.

- Zero-disagreement matches are the only ones where nothing in the dataset pushes back at all — still not a guarantee, just the cleanest read available.

- This document deliberately isolates the opposing case only. Pair it with a separate confirming-evidence pass (component-by-component P1 vs P2 differential, contradiction check, final Green/Yellow/Red) rather than using disagreement count alone as the verdict.

- State plainly if either the independent Green-list cross-reference source or full external (live) validation wasn't available/run for a given batch — don't drop those steps silently.

PART 2 — PREDICTED-LOSER / TRAP-GAME DETECTION (same extensiveness as the original verification audit, run in the opposite direction)

Purpose: Part 1 asks "what disagrees with the favorite." Part 2 builds the strongest honest case *for the underdog* — the player the Matrix predicts to lose — using the same full metric-family sweep as a full verification audit, just pointed at the other player. A "trap game" is a match where the underdog's case is unusually strong even though the Matrix still favors the other player.

Same rule as Part 1: nothing invented. Every line traces to a real field. Mark UNAVAILABLE where the data doesn't exist rather than guessing.

2.1 — What the backtest actually shows about upsets (recompute this block from the current verified workbook — these are reusable *categories*, the numbers must be refreshed each time)

Historical mining of the verified-results workbook should always include:

| Question | How to compute | What it's for |

|---|---|---|

| Upset rate by Closeness tier | losses in tier / total in tier | Which closeness label most often hides a live underdog |

| Upset rate by Recommendation label | losses in label / total in label | Which confidence label is actually least trustworthy |

| Upset rate by WP band | losses in band / total in band | Which probability range is most upset-prone |

| **Self-flagged Upset Risk vs actual upsets** | % of actual upsets that were flagged "Low" upset risk by the Matrix itself | Tests whether the Matrix's own upset warning is worth anything |

| Data Quality: losses vs wins (mean) | compare average DQ between the two groups | Tests whether thin data actually predicts upsets, or is a red herring |

| Model Agreement: losses vs wins | upset rate within "Strong" vs "Mixed" vs "High Disagreement" | Tests whether stated agreement protects against upsets |

| Monte Carlo support < 55% for the favorite | % of upsets vs % of correct picks where MC support was under 55% | A live, checkable early-warning threshold |

| Monte Carlo support < 50% for the favorite | same, at the 50% line (simulation numerically disagrees with headline) | Flags matches in territory the backtest sample may not even contain — treat as *more* uncertain than the sample average, not less |

Page 4

| Headline-vs-MC gap size, losses vs wins | mean/median gap in each group | Tests whether a big headline-vs-simulation gap alone is diagnostic (in the one sample run, it was NOT strongly diagnostic on its own — worth re-testing on fresh data, not assuming) |

| Straight-set probability gap, losses vs wins | compare predicted-winner vs opponent straight-set % |

Weak/inconclusive in the one sample tested — don't over-weight this alone |

Findings from the sample run on file (100 matches, 30 upsets) — treat as a prior to test against future data, not a permanent law:

- The Matrix's own "Upset Risk" flag caught almost none of the actual upsets — 29 of 30 upsets were labeled "Low" upset risk

beforehand. Do not use the Matrix's self-flag alone to rule a match safe from being a trap game.

- Data Quality was statistically identical between upsets and correct picks (~80% either way). Thin data is not what predicts upsets in this sample — don't lean on DQ as a trap-game filter.
- "Moderate lean" produced the *most* upsets in raw count (57% of all upsets) even though "Close matchup" had the *highest rate* (50% vs 27%). A trap-game scan has to cover the full "Moderate lean" pool, not just matches literally labeled a coin flip — most upsets by volume come from there simply because there are more of those matches.
- "High Confidence" recommendation carried the single highest upset rate of any label (41.7%) — higher than "Moderate Confidence" (30%) and "Elite" (27%). Treat "High Confidence" as a caution flag, not reassurance, consistent with Part 1.
- Monte Carlo support under 55% for the favorite showed up in 26.7% of upsets vs 18.6% of correct picks — a real but modest lean; not a standalone trigger by itself.
- Monte Carlo never actually dropped below 50% for the stated favorite anywhere in the sample tested. Any live match where it does is outside the historical envelope — flag it as higher-uncertainty than anything the backtest can vouch for, not lower.

2.2 — Underdog case-building: full metric-family sweep

For every match, run the underdog through the same families a full verification audit would run the favorite through. Mark each FOR (helps the underdog), AGAINST (helps the favorite), or UNAVAILABLE.

Matrix Headline & Confidence

- Headline WP margin — how close to 50% is it really?
- Recommendation label — is it one of the historically weak labels (see 2.1)?
- Upset Risk as self-flagged — note it, but weight it low per the finding above.
- Model Agreement label — "Strong" does not reliably rule out an upset (see 2.1); check anyway for what it's worth.
- Data Quality — note it, but don't treat low DQ as the reason to expect an upset, or high DQ as protection.

Surface / Elo

- Underdog's Surface Elo raw probability and reliability/sample — is it actually close to the favorite's once you set aside the blend weighting?
- Underdog's Elo trend — improving, stable, or declining independent of the raw number.
- Surface specialization label — is the underdog the surface specialist here even though they're not the overall favorite?

Serve / Return

- Underdog's Serve/Return raw probability and reliability.
- Does the underdog's return profile specifically attack the favorite's serve profile (style/matchup fit), separate from the raw number?

Recent Form

- Underdog's Recent Form raw probability, trend direction, and reliability.
- Any explicit "improving" trend flag on the underdog, especially paired with a "declining" or "stable" flag on the favorite.
- Opponent-adjusted form for the underdog, if present — are their recent results against tougher or weaker competition than the raw form score implies?

Page 5

H2H

- Raw H2H number and reliability/sample. In most Matrix exports seen so far this is near-zero and excluded — say so plainly rather than treating a 50/50 H2H number as a real signal either way.

Market

- Does the market component favor the underdog, or at least sit closer to 50/50 than the headline WP? A market that disagrees with (or is much less confident than) the model's own headline is one of the more credible trap-game signals, since market data aggregates outside information the Matrix's other components may not have.

General / Specialist Model

- Underdog's raw probability on each, and reliability/sample/status (Active vs Excluded, Low vs High availability). A specialist model favoring the underdog on a thin sample is weaker evidence than the same signal on a deep one — say which it is.

Fatigue / Physical

- Underdog's fatigue index vs the favorite's. A fresher underdog against a fatigued favorite is one of the more concrete, checkable trap-game signals available in this data.
- Matches played in the trailing window for both, if present.

Monte Carlo / Set Score

- Underdog's implied probability (100 – favorite's MC%).
- Is the favorite's MC support under 55% (modest historical upset lean) or under 50% (numerically favors the underdog, outside the historical sample entirely — flag as high-uncertainty)?
- How far does the simulation range extend in the underdog's favor?
- Does the predicted set score (e.g., 2-0) actually match what the simulation supports, or is a 2-1/underdog-competitive outcome more consistent with the range?

Style / Matchup

- Style labels for both players (e.g., all-court, clay specialist, big server) — does the underdog's style specifically counter the favorite's, independent of the surface Elo number?

Ranking / External

- UNAVAILABLE unless a dedicated external-validation pass has actually been run (see Part 1's completion-check note) — do not infer ranking trajectory or current form from Matrix fields that don't contain it.

2.3 — Trap Game Score (composite)

For each match, count how many of the following are true for the underdog. This is a simple point count, not a probability — it exists to surface matches worth a second look, not to replace the audit above.

- ☐ Underdog favored by at least one live model component (Elo, S/R, Form, Market, General, or Specialist)
- ☐ Underdog favored by two or more live model components
- ☐ Market specifically favors the underdog, or sits materially closer to 50/50 than the headline WP
- ☐ Favorite's Monte Carlo support is under 55%
- ☐ Favorite's Monte Carlo support is under 50% (outside the historical sample envelope entirely)
- ☐ Underdog has a better (lower) fatigue index than the favorite
- ☐ Underdog's Recent Form trend is flagged "improving" while favorite's is "stable" or "declining"
- ☐ Match is labeled "Moderate lean" or closer (not "Clear favorite") — this is the pool where most historical upsets actually live
- ☐ Recommendation label is "High Confidence" specifically (the weakest-performing label in the one sample tested)

Page 6

- ☐ Style/specialization label favors the underdog on this surface

Scoring guide (recalibrate against a larger sample if the pattern doesn't hold):

- 0–1 checked: no meaningful trap-game case — the audit found the underdog case, but it's thin.
- 2–4 checked: Trap Game Watch — worth a second look before treating the favorite as safe.
- 5+ checked: High Trap Risk — the underdog's case is comparable in breadth to what Part 1 flags as a heavily-disagreed-with favorite.

Treat the favorite pick with real skepticism.

2.4 — Per-match underdog entry template (copy for every matchup, alongside the Part 1 block for the same match)

[Underdog] — case as predicted loser vs [Favorite] ([favorite's headline WP]%)

Trap Game Score: [X] / 10 checked → [No case / Trap Game Watch / High Trap Risk]

- Component support: [list any of Elo/S-R/Form/Market/General/Specialist favoring the underdog, with contribution and reliability]
- Market read: [favors underdog / closer to 50-50 than headline / unavailable]
- Monte Carlo: favorite supported at [X]% ([above 55% - no lean / 50-55% - modest lean / below 50% - outside historical sample, flag high]); simulation range extends to [low]% for favorite, i.e. up to [100-low]% for underdog in some simulated outcomes
- Fatigue: underdog [fresher/more fatigued] than favorite ([underdog FI] vs [favorite FI])

- Form trend: underdog [improving/stable/declining], favorite [improving/stable/declining]
- Closeness/label context: labeled '[closeness]', recommendation '[label]' — [note if this is one of the historically weak/upset-prone labels from 2.1]
- Style/surface fit: [note if underdog's style or specialization counters the favorite]
- Ranking/external: UNAVAILABLE unless external validation was run for this match

Verdict: [No meaningful trap-game case / Worth a second look / High trap risk — treat favorite skeptically]

PART 3 — FULL PREDICTED-LOSER VERIFICATION AUDIT (Mirror of the

Winner Verification Audit)

Purpose: This is the anti-favorite version of the full verification audit. The goal is not merely to find disagreement with the Matrix pick.

The goal is to determine whether the player predicted to lose has a *real, repeatable, opponent-specific pathway to winning* that the favorite-side model may be underweighting.

For every matchup, treat the predicted loser as a candidate winner and run the entire verification framework in reverse. Do not stop after finding one attractive underdog statistic. The underdog must survive the same reliability, sample-size, opponent-adjustment, historical-validation, and external-validation standards as a projected winner.

3.1 — Mandatory evidence classification

Every item evaluated for the predicted loser must receive one of these labels:

- UD+ STRONG: materially supports the predicted loser, with reliable sample/relevance.
- UD+ WEAK: supports the predicted loser but sample/reliability/relevance is limited.
- FAV+ STRONG: materially supports the favorite and directly weakens the trap-game case.
- FAV+ WEAK: mildly supports the favorite.
- CONFLICT: credible evidence exists on both sides.

Page 7

- UNAVAILABLE: no trustworthy data supplied/found.
 - EXCLUDED: field exists but should carry zero decision weight because the Matrix itself excludes it or reliability is effectively unusable.
- Do not count UD+ WEAK the same as UD+ STRONG. Do not count five correlated metrics as five independent signals.

3.2 — Full underdog metric families

A. Headline / Calibration Vulnerability

Evaluate:

- Favorite headline WP and underdog implied WP.
- Distance from 50/50.
- Headline probability vs Monte Carlo probability.
- Headline probability vs market no-vig probability.
- Headline probability vs General model.
- Headline probability vs Specialist model.
- Headline probability vs Surface Elo probability.
- Headline probability vs Serve/Return probability.
- Headline probability vs Recent Form probability.
- Maximum component disagreement spread.
- Whether the headline appears to be carried by one heavily weighted model.
- Whether the confidence label historically overstates actual win rate.
- Whether this probability bucket is overconfident in the verified backtest.
- Favorite calibration error by price/probability bucket, if enough sample exists.

Trap signal: headline says "safe favorite" while independent components cluster materially closer to 50/50.

B. Surface Strength / Surface Mispricing

Evaluate both players:

- Surface Elo.
- Elo differential.
- Surface Elo probability.
- Surface sample depth.
- Effective weighted sample.
- Surface Elo reliability.
- Surface Elo trend/momentum.
- Last-52-week surface record.
- Current surface swing.
- Surface-transition performance.
- Fast/medium/slow hard split where available.
- Indoor/outdoor split where available.
- Tournament-specific surface history.
- Court-speed fit.
- Ball-type fit.

Page 8

- Altitude/weather interaction when applicable.
- Style-adjusted Elo.

Trap signal: favorite owns the higher generic Elo, but underdog has the more relevant current-surface, court-speed, or style-adjusted profile.

C. Serve Profile — Favorite Vulnerability vs Underdog Attack

Evaluate:

- Hold %.
- Service points won %.
- First-serve-in %.
- First-serve points won %.
- Second-serve points won %.
- Ace rate.
- Double-fault rate.
- Break points saved %.
- Clean-hold rate.
- Hold vulnerability.
- Serve-out reliability.
- Serve speed/first-serve decay by set.
- Second-serve deterioration.
- Serve dependency index.
- Serve placement predictability when available.
- End-of-set serve deterioration.

Then ask specifically:

- Can the underdog attack the favorite's second serve?
- Does the favorite rely too heavily on first-serve dominance?
- Is the favorite vulnerable when serve quality drops even slightly?

Trap signal: favorite's entire advantage depends on a serve level the underdog is specifically equipped to neutralize.

D. Return Profile — Underdog Upset Weapon

Evaluate:

- Break %.
- Return points won %.
- First-serve return points won.
- Second-serve return points won.
- Break points created per return game.
- Break-point conversion %.
- Return pressure rate.
- Clean-break rate.
- Return-game threat consistency.
- Return +1 effectiveness.

Page 9

- Second-serve attack rate.
- Return position compatibility.
- Return-quality trend.
- Return dependency index.

Trap signal: underdog owns the strongest opponent-specific weapon in the matchup: a return profile that attacks the favorite's known serve weakness.

E. Serve-vs-Return Interaction

Compute or assess:

- Favorite expected hold vs this underdog.
- Underdog expected hold vs this favorite.
- Favorite expected break vs this underdog.
- Underdog expected break vs this favorite.
- Expected hold/break differential.
- Serve-vs-return compatibility score.
- Return-vs-serve compatibility score.
- Second-serve mismatch index.
- Serve–return interaction residual.
- Opponent-specific break expectancy.
- Opponent-specific hold expectancy.

Trap signal: raw serve/return averages favor the favorite, but the matchup-specific interaction favors the underdog.

F. Recent Form — Real vs Fake Form

For both players evaluate:

- Last 5.
- Last 10.
- Current streak.
- Form score.
- Trend direction.
- Performance acceleration.
- Quality of last 3–5 matches.
- Opponent-adjusted recent form.
- Opponent-quality-adjusted record trend.
- Hidden improvement despite losses.
- Hidden decline despite wins.
- Observed vs expected W/L.

- Quality win rate.
- Bad-loss rate.
- Bad-loss severity.
- Blowout-loss frequency.
- Post-breakthrough performance.

Page 10

- Post-upset regression.
- Post-title/final hangover.
- Post-long-run fatigue.
- Grand Slam hangover.

Trap signal: favorite looks better in raw W/L, but underdog is improving against stronger opposition while favorite's recent wins are low-quality or unsustainable.

G. Loss Autopsy — How the Favorite Actually Loses

For the favorite's recent losses evaluate:

- Favorite status in each loss.
- Opponent quality.
- Surface.
- Point differential.
- Break differential.
- Serve deterioration.
- Return deterioration.
- Lost after leading.
- Lost Set 1.
- Deciding-set loss.
- Tiebreak loss.
- Physical issue.
- Match length.
- Competitive vs blowout.
- Bad-loss severity.

Then compare today's underdog to those opponents:

- Similar Elo?
- Similar serve style?
- Similar return quality?
- Similar rally profile?
- Similar handedness?
- Similar court preference?

Trap signal: today's underdog resembles the exact player archetype that repeatedly causes the favorite to fail.

H. Underdog Win Autopsy — Are the Upsets Repeatable?

For the underdog's strongest recent wins evaluate:

- Was the opponent actually favored?
- Surface relevance.
- Opponent Elo/ranking.
- Did underdog win through serve dominance, return pressure, long rallies, tiebreak variance, or opponent collapse?
- Was the win point-dominant or clutch-dependent?
- Was it retirement-distorted?

- Was it a one-off outlier or part of a repeatable pattern?
- Quality win rate.
- Upset-quality score.

Trap signal: underdog has recently beaten players similar to today's favorite using a repeatable pathway, not random variance.

I. Common Opponents / Network Evidence

Evaluate:

- Direct common opponents.
- Divergent outcomes.
- Scoreline comparison.
- Same-surface common opponents.
- Same-level common opponents.
- Recency-weighted common-opponent results.
- Opponent-strength weighting.
- Common-opponent point differential.
- Hold differential.
- Break differential.
- Straight-set differential.
- Set-1 differential.
- Second-degree opponent network.
- Transitive performance score.
- Network Elo.

Trap signal: underdog consistently performs better than favorite against shared or highly comparable opposition.

J. H2H / Rematch Dynamics

Evaluate:

- Overall H2H.
- Same-surface H2H.
- Recency.
- Scoreline.
- Match competitiveness.
- Junior/ITF H2H when relevant.
- Tactical pattern from previous meeting.
- Rematch adjustment.
- Revenge/rematch tactical differential.
- Whether prior loser corrected the exact weakness exposed previously.

Trap signal: underdog has demonstrated an opponent-specific tactical solution that generic model components fail to capture.

K. Set-1 / Slow-Start Trap Risk

Evaluate both players:

- Set-1 win rate.

- First-break frequency.
- Opening service-game hold %.
- Opening return-game break %.
- First-four-games differential.
- First-six-games point differential.

- Early break-conceded frequency.
- Time to first break.
- Set-1 slow-start index.
- First-set recovery after early break.
- Warm-up dependency.

Trap signal: favorite is a chronic slow starter while underdog frequently breaks first and protects early leads.

L. Comeback / Closing / Lead Durability

Evaluate:

- Record after losing Set 1.
- Record after winning Set 1.
- Second-set performance after losing Set 1.
- Deciding-set win rate.
- Break-back rate.
- Break-consolidation rate.
- Serving-for-set conversion.
- Serving-for-match conversion.
- Lead durability.
- Deficit survivability.
- Double-break creation/surrender.
- Rebreak-window probability.
- Opponent one-set-up collapse rate.
- Opponent serving-for-match failure rate.

Trap signal: favorite often gets into winning positions but fails to close, while underdog is unusually resilient from behind.

M. Pressure / Tiebreak / Clutch Sustainability

Evaluate:

- Pressure-point performance.
- Clutch hold/break.
- Tiebreak record.
- Expected vs actual tiebreak wins.
- Expected vs actual deciding-set wins.
- Deuce-game win %.
- Break-point quality.
- Break opportunities needed per break.
- Multiple-BP survival/conversion.

Page 13

- Extended-deuce endurance.
- Tiebreak entry quality.
- Mini-break recovery/consolidation.

Trap signal: favorite's record depends heavily on unsustainably good clutch outcomes, while underdog's underlying point dominance is stronger than results indicate.

N. Straight-Set / Match-Length Vulnerability

Evaluate:

- Favorite straight-set win rate.
- Underdog straight-set loss rate.
- Straight-set rate vs comparable opposition.

- MC straight-set probability for each player.
- Matrix 2-0 call vs actual simulation support.
- Probability match reaches Set 3.
- Favorite performance when forced to Set 3.
- Underdog decider performance.
- Long-match resilience.
- Physical-cliff probability.

Trap signal: Matrix predicts 2-0, but matchup is structurally likely to reach a decider where underdog becomes much more dangerous.

O. Fatigue / Scheduling / Recovery

Evaluate:

- Fatigue index.
- Matches in 7/14/30 days.
- Minutes on court.
- Sets/games played.
- Recent three-setters.
- Late finishes.
- Rest hours.
- Qualifying workload.
- Consecutive-day play.
- Consecutive tournament count.
- Travel/time zones.
- Recovery hours.
- Back-to-back long-match penalty.
- Tournament adaptation advantage.
- Qualifier adaptation bonus.
- Accumulated workload debt.
- Recovery efficiency.
- Circadian mismatch.

Trap signal: favorite is stronger on paper but enters with a material physical/recovery disadvantage that the underdog can exploit.

Page 14

P. Injury / Availability / Physical Decline

Evaluate:

- Injuries.
- Recent withdrawals.
- Retirements.
- Medical timeouts.
- Illness.
- Layoff.
- Return-from-injury form.
- Match fitness.
- Movement asymmetry.
- Serve velocity decline.
- First-serve decline.
- Return-quality decline.
- Late-match physical decline.

Trap signal: favorite's nominal rating is stale because current physical capacity is below the level represented by Elo/ranking.

Q. Ranking / Development / Mispricing

Evaluate:

- Current ranking.
- Ranking trajectory.
- Rapid riser/faller.
- Ranking-performance disconnect.
- Age/development curve.
- Ranking catch-up estimate.
- Elo catch-up estimate.
- Protected-ranking status.
- Tour-level transition performance.
- Post-breakthrough sustainability.

Trap signal: underdog is a rapidly improving player whose ranking/Elo has not caught up to current level.

R. Style / Rally / Tactical Collision

Evaluate where available:

- Rally-length profile.
- First-strike efficiency.
- Neutral-rally efficiency.
- Long-rally tolerance.
- Forehand/backhand effectiveness.
- Weak-wing exploitation.
- Movement exploitation.
- Net-point frequency/success.

Page 15

- Passing-shot success.
- High-ball/low-ball tolerance.
- Slice exposure.
- Lefty/righty interaction.
- One-handed vs two-handed backhand interaction.
- Serve-placement matchup.
- Return-position compatibility.
- Preferred-pattern collision.
- Plan-B effectiveness.
- Matchup adaptability.

Trap signal: favorite's preferred pattern feeds directly into underdog's best tactical pattern.

S. Market Trap Detection

Evaluate:

- Multi-book moneyline consensus.
- No-vig probability.
- Opening price.
- Current price.
- Closing price if available.
- Direction/magnitude of line movement.
- Sharp vs recreational book divergence.

- Market disagreement dispersion.
- Prediction-market consensus.
- Late-line acceleration.
- Model-vs-market divergence.
- Player-specific price calibration.
- Favorite performance by price bucket.

Flag especially:

- Matrix 70%+ but market only ~55–60%.
- Favorite drifting materially despite positive public-facing model signals.
- Underdog shortening across sharper books.
- Books disagree sharply with one another.

Trap signal: independent market information persistently prices the underdog stronger than the Matrix does.

T. Favorite Fragility / Underdog Upset Compatibility

Evaluate:

- Favorite Fragility Index.
- Upset resistance.
- Favorite failure-mode profile.
- Opponent failure-mode compatibility.
- Underdog upset compatibility score.

Page 16

- Favorite performance when opponent holds first 3 service games.
- Favorite after failing early break chances.
- Favorite after losing first break.
- Favorite at 4-4.
- Favorite in tiebreak sets.
- Favorite when forced to Set 3.

Trap signal: the underdog's strengths map directly onto the favorite's documented failure conditions.

U. Motivation / Context / Venue

Use only documented, reliable information:

- Points-defending pressure.
- Seeding implications.
- Career milestone pressure.
- Home-country effect.
- Crowd support.
- Tournament-specific history.
- Prior-year round reached.
- Short-notice entry.
- Lucky-loser status.
- Wildcard effect.
- Walkover rhythm effect.
- Coaching changes.
- Equipment changes.
- Travel friction.
- Weather sensitivity.
- Venue familiarity.

Do not invent psychology. Only count documented context when evidence exists.

PART 4 — FAVORITE FAILURE-MODE AUDIT

For every favorite, answer these questions explicitly.

4.1 — What must go right for the favorite to win?

List the favorite's top 3–5 dependencies:

- first serve must dominate?
- must break early?
- must keep rallies short?
- must avoid Set 3?
- must win tiebreaks?
- must prevent underdog attacking second serve?
- must maintain movement/endurance?

Page 17

4.2 — What causes this favorite to lose?

Use recent-loss autopsy and historical tendencies to identify 3–5 failure modes.

4.3 — Can today's underdog actually create those conditions?

For each favorite failure mode assign:

- YES — strong evidence
- POSSIBLY — partial evidence
- NO — underdog lacks the weapon
- UNAVAILABLE

4.4 — Failure-Mode Compatibility Count

Count the number of favorite failure modes today's underdog can realistically trigger.

- 0: low trap compatibility.
- 1: mild.
- 2: meaningful.
- 3+: high trap compatibility.

This is more important than simply counting how many raw metrics disagree with the favorite.

PART 5 — UNDERDOG WIN-PATHWAY AUDIT

For every predicted loser, assess each distinct pathway to victory.

1. Serve-through pathway — can underdog hold often enough to force high-variance sets/tiebreaks?
2. Return-pressure pathway — can underdog repeatedly attack favorite's serve?
3. Second-serve exploitation pathway.
4. Short-rally pathway.
5. Long-rally pathway.
6. Movement/physical pathway.
7. Slow-start pathway — favorite gives away Set 1 often.
8. Deciding-set pathway.
9. Tiebreak pathway.
10. Fatigue pathway.
11. Style mismatch pathway.
12. Market-information pathway.
13. Favorite-collapse pathway.
14. Surface-transition pathway.
15. Recent-improvement / ranking-lag pathway.

For each pathway label:

- REALISTIC

- WEAK

Page 18

- NOT SUPPORTED

- UNAVAILABLE

Then report Underdog Win Pathways Count = number of REALISTIC distinct pathways.

Important: two pathways built on the same underlying fact should not be counted as independent.

PART 6 — INDEPENDENT EVIDENCE / CIRCULARITY AUDIT

A trap can be missed when five model outputs all agree because they recycle the same underlying data.

For every matchup:

- Count raw agreeing components.
- Estimate genuinely independent evidence families.
- Identify likely overlap between General and Specialist models.
- Identify overlap between Surface Elo and opponent-strength adjustments.
- Identify overlap between Recent Form and Specialist/General inputs.
- Identify market as an outside-information family when truly independent.
- Identify external injury/news as another independent family.

Report:

- Raw Favorite Agreement Count: X
- Independent Favorite Evidence Count: Y
- Raw Underdog Agreement Count: X
- Independent Underdog Evidence Count: Y
- Evidence Independence Risk: Low / Moderate / High

Trap signal: favorite appears to have 5–6 supporting components but only 2 genuinely independent evidence families.

PART 7 — FAVORITE STRESS TEST / WINNER-SWITCH THRESHOLD

For every favorite, deliberately worsen the favorite's assumptions within realistic bounds:

- reduce favorite Elo edge modestly.
- reduce favorite recent-form edge.
- reduce favorite serve strength.
- improve underdog return strength.
- increase favorite fatigue.
- move market toward underdog.
- remove the single strongest model component.
- remove General Model.
- remove Specialist Model.
- remove market.
- remove recent form.

Page 19

- remove Elo.

Then ask:

- Does the favorite still remain the predicted winner?
- How many single-component removals flip the winner?
- What is the smallest realistic input change that flips the pick?

Report:

- Favorite Robustness: Strong / Moderate / Fragile.

- Winner-Switch Threshold: High / Medium / Low.

- Counterfactual Survival: X of Y removals.

Trap signal: favorite flips under one modest assumption change or after removing one dominant component.

PART 8 — HISTORICAL TRAP-TWIN SEARCH

Search the verified workbook for matches most similar to the current favorite/underdog profile using as many available dimensions as possible:

- headline WP.

- closeness.

- recommendation.

- upset risk.

- DQ.

- model agreement.

- Elo gap.

- serve/return gap.

- recent-form gap.

- market direction.

- active disagreement count.

- MC support.

- headline-vs-MC gap.

- straight-set gap.

- fatigue gap.

- opponent-adjusted form.

For every current match identify:

- Top positive twins where favorite won.

- Top negative twins where favorite lost.

- Similarity score or explanation.

- Number of comparable historical matches.

- Favorite win rate within those comparable matches.

Do not call a historical pattern strong if fewer than ~10 genuinely comparable examples exist.

Trap signal: nearest historical twins contain a disproportionately high favorite-loss rate.

Page 20

PART 9 — EXTERNAL TRAP VALIDATION

Only after the internal disagreement/underdog audit, verify both players externally.

At minimum check:

- current ranking.

- ranking movement.

- last 5 results.

- last 10 results.

- same-surface record.

- current tournament results.

- current field/draw validity.

- H2H.

- injuries.

- withdrawals.

- recent retirements.
- medical concerns.
- rest/fatigue context.
- travel.
- recent notable wins/losses.
- current sportsbook market from multiple books when available.
- current prediction-market information when available.

For lower-tour players, check multiple reliable result/stat sources where possible because identity and stale data issues are more common.

External evidence must be classified:

- CONFIRMS TRAP
- WEAKLY CONFIRMS TRAP
- NEUTRAL
- WEAKENS TRAP
- INVALIDATES TRAP

Never use external validation to cherry-pick only the underdog case. Search for evidence that strengthens the favorite too.

PART 10 — ENHANCED TRAP GAME SCORE

Do not rely on one raw count. Build four sub-scores.

A. Internal Disagreement Score (0–10)

Points for meaningful, reliable disagreement:

- +2 market favors underdog.

Page 21

- +2 two or more reliable live components favor underdog.
- +1 one reliable live component favors underdog.
- +1 favorite MC <55%.
- +2 favorite MC <50%.
- +1 headline-MC gap ≥ 15 .
- +1 favorite is more fatigued.
- +1 underdog recent-form trend better.
- +1 meaningful H2H/same-surface H2H advantage with real sample.
- +1 style/surface matchup favors underdog.

Cap at 10.

B. Favorite Fragility Score (0–10)

- +2 favorite depends heavily on one component.
- +2 winner flips when strongest component removed.
- +1 poor historical calibration bucket.
- +1 close/moderate-lean tier.
- +1 favorite has recent physical concern.
- +1 favorite slow-start risk aligns with underdog strength.
- +1 favorite weak deciding-set/closing profile.
- +1 favorite has recent bad-loss pattern vs similar opponents.
- +1 favorite's current ranking/Elo looks stale vs current form.

Cap at 10.

C. Underdog Win-Pathway Score (0–10)

+1 for each genuinely independent REALISTIC pathway from Part 5, cap at 10.

D. Historical / External Confirmation Score (0–10)

- +2 negative historical twins strongly resemble current match.
- +1 current backtest bucket is upset-prone.
- +2 external market confirms underdog strength.
- +2 current form/external results confirm underdog.
- +1 H2H confirms underdog.
- +1 physical/fatigue evidence confirms underdog.
- +1 current-surface evidence confirms underdog.

Cap at 10.

Composite Trap Risk = A + B + C + D (0–40)

Suggested interpretation — recalibrate using future results:

- 0–7: LOW TRAP RISK
- 8–14: WATCH
- 15–21: MODERATE TRAP

Page 22

- 22–28: HIGH TRAP RISK
- 29–40: EXTREME TRAP / FAVORITE SHOULD BE RE-AUDITED BEFORE USE

Never convert this score directly into a win probability unless separately calibrated on a large enough sample.

PART 11 — TRAP-SPECIFIC RED FLAGS THAT OVERRIDE A GREEN LABEL

A favorite cannot remain "clean Green" without explicit review if any of these occur:

1. Market favors the underdog while Matrix headline is $\geq 65\%$ favorite.
2. Favorite MC support $< 55\%$.
3. Favorite MC $< 50\%$.
4. Headline-MC gap ≥ 20 points.
5. Two or more high-reliability independent components favor underdog.
6. Favorite's biggest supporting component is low-sample/low-reliability.
7. Favorite has active injury/return-from-injury concern.
8. Underdog has strong same-surface H2H with adequate sample.
9. Underdog has better current form after opponent adjustment.
10. Favorite recently lost repeatedly to the same archetype as today's underdog.
11. Favorite is materially more fatigued and underdog has a physical/long-match pathway.
12. External market moves materially toward underdog.
13. Match is in historically weak favorite bucket *and* underdog owns at least two independent pathways.
14. Favorite pick flips under one realistic stress-test adjustment.
15. Historical nearest-neighbor matches show materially elevated upset rate.

Any override flag requires a written explanation before assigning Green.

PART 12 — FINAL PER-MATCH TRAP AUDIT OUTPUT

For every matchup, output this block:

[Favorite] vs [Underdog]

Matrix favorite: [name] — [headline WP]%

Predicted loser / underdog: [name]

Current Matrix color: [if supplied]

1. Favorite headline/calibration:

- Headline WP:
- MC:

- Market no-vig:
- Headline-MC gap:
- Headline-market gap:
- Historical calibration bucket:

Page 23

2. Predicted loser evidence:

- Surface/Elo:
- Serve:
- Return:
- Recent form:
- Opponent adjustment:
- Common opponents:
- H2H/rematch:
- Pressure/closing:
- Fatigue/physical:
- Ranking/development:
- Style/matchup:
- Market:
- Monte Carlo/set-score:

3. Reliability / sample check:

- Strong underdog evidence:
- Weak underdog evidence:
- Excluded evidence:
- Missing/unavailable evidence:

4. Favorite contradictions:

- [all meaningful contradictions]

5. Favorite failure modes:

- Failure mode 1 → can underdog trigger? [Yes/Possibly/No]
- Failure mode 2 → ...
- Failure mode 3 → ...

6. Underdog realistic win pathways:

- [pathway + evidence]
- Total independent realistic pathways: [N]

7. Circularity / evidence independence:

- Raw favorite agreement:
- Independent favorite evidence:
- Raw underdog agreement:
- Independent underdog evidence:

8. Favorite stress test:

- Robustness:
- Winner-switch threshold:
- Counterfactual survival:

9. Historical trap twins:

Page 24

- Positive favorite twins:

- Negative/upset twins:
 - Comparable sample:
 - Historical favorite win rate:
10. External validation:
- Ranking/current form:
 - Same surface:
 - Injury/availability:
 - H2H:
 - Market:
 - Current draw validity:
 - External verdict: [confirms/neutral/weakens/invalidates trap]
11. Trap scoring:
- Internal Disagreement: [x/10]
 - Favorite Fragility: [x/10]
 - Underdog Win Pathways: [x/10]
 - Historical/External Confirmation: [x/10]
 - Composite Trap Risk: [x/40]
12. Final trap verdict:
- LOW TRAP / WATCH / MODERATE / HIGH / EXTREME
 - Favorite status: KEEP / DOWNGRADE / REMOVE / FLIP / PASS
 - Evidence winner:
 - Reason:

PART 13 — BATCH-LEVEL OUTPUT

After all matches are audited, produce:

A. Predicted Losers Most Likely to Win

Rank underdogs from strongest upset case to weakest.

| Rank | Predicted Loser | Favorite | Composite Trap Risk | Independent Win Pathways | Market Conflict |

Historical Twin Risk | Final Action |

|---|---|---|---|---|---|

B. Favorite Trap Games

Rank favorite picks most likely to fail.

C. False Greens

List any Green favorite that triggered one or more override flags.

D. False Confidence Labels

Page 25

List "Elite," "High Confidence," or "Strong Agreement" favorites whose underlying evidence is materially contested.

E. Matrix Picks to Flip

Only include flips where underdog has stronger *reliable* evidence overall.

F. Pass / No-Edge Matches

List matches where neither side earns a reliable edge.

G. Safest Favorites After Trap Audit

Favorites that survived every underdog/failure-mode/stress-test layer.

PART 14 — COMPLETION CHECK — NO SHORTCUTS

Before claiming the disagreement/trap audit is complete, verify:

[] Entire verified backtest workbook recalculated from actual rows.

- ☐ All current matchups extracted and deduplicated.
- ☐ Favorite evaluated separately.
- ☐ Predicted loser evaluated separately.
- ☐ Entire master metrics list screened for both players.
- ☐ Missing metrics explicitly marked unavailable rather than neutral.
- ☐ Reliability/sample/status checked for every model component.
- ☐ General/Specialist overlap/circularity considered.
- ☐ Direct favorite-vs-underdog differential completed.
- ☐ Favorite loss-autopsy completed where data exists.
- ☐ Underdog win-autopsy completed where data exists.
- ☐ Common-opponent/network layer completed where available.
- ☐ Set-1/slow-start layer completed where available.
- ☐ Deciding-set/closing layer completed where available.
- ☐ Fatigue/physical layer completed.
- ☐ Favorite failure modes identified.
- ☐ Underdog win pathways identified.
- ☐ Favorite stress test performed.
- ☐ Winner-switch threshold assessed.
- ☐ Historical trap twins searched.
- ☐ External validation run for both players.
- ☐ Current market checked from multiple sources where available.
- ☐ Current draw/withdrawal validity checked.
- ☐ Composite Trap Risk built from documented inputs only.
- ☐ Green override flags checked.
- ☐ Final KEEP/DOWNGRADE/REMOVE/FLIP/PASS assigned only after every prior step.

If any required source is missing, state exactly which layer is incomplete. Do not say "complete" if it is not.

Page 26

PART 15 — PRIMARY OPERATING PRINCIPLE

The disagreement audit is not finished when it proves that the favorite has weaknesses. It is finished only when it answers both questions:

1. Why can the favorite lose?
2. Does this specific underdog possess the actual weapons, form, physical condition, matchup fit, market support, and historical profile required to make those failure modes happen today?

A true trap game requires both sides of that equation. The goal is to detect the predicted loser *before the upset happens*, without turning every imperfect favorite into a fake trap signal.

PART 16 — NON-NEGOTIABLE EXECUTION ENFORCEMENT

This section is mandatory and exists specifically to prevent incomplete or shortcut audits from being presented as complete.

Before producing ANY final trap verdict, the assistant must create an internal completion ledger for every matchup.

For each matchup, the ledger must contain:

1. Favorite-side metric sweep completed: YES/NO
2. Underdog-side metric sweep completed: YES/NO
3. Reliability/sample/status reviewed: YES/NO
4. Direct differential completed: YES/NO
5. Favorite loss-autopsy completed or explicitly UNAVAILABLE: YES/NO
6. Underdog win-autopsy completed or explicitly UNAVAILABLE: YES/NO

7. Failure-mode compatibility completed: YES/NO
8. Independent win-pathways counted: YES/NO
9. Circularity/independent-evidence review completed: YES/NO
10. Stress test performed: YES/NO
11. Winner-switch threshold assessed: YES/NO
12. Historical trap twins searched: YES/NO
13. External validation completed for BOTH players: YES/NO
14. Current market checked from multiple sources where available: YES/NO
15. Draw/withdrawal validity checked: YES/NO
16. Green override rules checked: YES/NO
17. Composite Trap Risk calculated from documented evidence only: YES/NO

HARD STOP RULE

If ANY required ledger item is NO, the assistant may NOT:

- call the audit complete,
- label a match a confirmed trap,
- issue a final FLIP,
- issue a final safest-favorites list,

Page 27

- or claim all instructions were followed.

Instead, the assistant must say exactly:

INCOMPLETE AUDIT — missing: [specific layer/source].

NO COMPRESSION RULE

A batch summary may only be produced AFTER every matchup has its full per-match audit block completed.

NO INFERENCE SUBSTITUTION

Do not substitute:

- prior assistant verdicts,
- earlier Green lists,
- old market odds,
- old injury information,
- old rankings,
- remembered H2H,
- or generic player reputation

for current verification.

NO SILENT SKIPPING

Every unavailable requested metric must be explicitly classified UNAVAILABLE, EXCLUDED, or PARTIAL.

Missing data may never be treated as neutral, 50/50, or evidence for either player.

NO FAVORITE BIAS

The Matrix favorite begins the audit with zero presumption of correctness.

The predicted loser must be treated as a legitimate candidate winner until the full audit is complete.

NO TRAP BIAS

Do not manufacture a trap because the favorite has imperfections.

A trap requires BOTH:

- A) a documented favorite failure mode,

AND

- B) a documented, opponent-specific underdog pathway capable of triggering it.

FINAL CERTIFICATION

At the end of every completed batch audit, include:

AUDIT CERTIFICATION

- Matchups expected: [N]
- Matchups completed: [N]
- Favorite audits completed: [N/N]
- Underdog audits completed: [N/N]

Page 28

- Historical overlays completed: [N/N]
- External validations completed: [N/N]
- Stress tests completed: [N/N]
- Trap scores completed: [N/N]
- Missing required sources: [NONE or list]

Only if every required count equals N and no mandatory source is missing may the assistant write:

FULL DISAGREEMENT/TRAP AUDIT COMPLETE — NO REQUIRED STEP SKIPPED.

If any count is below N or any mandatory source is missing, the assistant must not use that completion statement.

PART 17 — ENFORCEMENT OVERRIDE PRECEDENCE

If any instruction elsewhere in this audit conflicts with Part 16, Part 16 controls.

The assistant must prioritize verifiability over speed, brevity, or convenience.

A shorter answer is never a valid reason to omit a required audit layer.

A lack of available data is never permission to guess.

A previous answer is never evidence unless the underlying source is independently re-checked in the current audit.

The goal is not merely to produce a trap list. The goal is to produce a trap list that can be traced back through every required evidence layer without hidden omissions.

PART 18 - FRESH ONLINE RESEARCH PROTOCOL - MANDATORY EVERY RUN

This section is a NON-NEGOTIABLE execution requirement.

Existing verification reports, previous assistant research, cached external validation, earlier Green audits, old web searches, and previously collected online information DO NOT satisfy the online-research requirement.

For EVERY matchup, perform a NEW online research cycle during the CURRENT disagreement/trap-audit run.

Research BOTH:

- A) the predicted favorite
- B) the predicted loser

The predicted loser must receive at least the same online-research depth as the favorite.

18.1 CURRENT MATCH VALIDITY

Freshly verify:

- tournament
- tournament level
- surface
- round

Page 29

- exact matchup
- both players remain in the draw
- withdrawal status
- replacement/lucky-loser status

- walkover status
- schedule/status when available

The exact matchup must be confirmed from a current reliable source whenever such a source is available.

If exact current draw validation cannot be obtained, mark:

CURRENT DRAW VALIDATION: INCOMPLETE

Do not silently rely on an older prediction sheet as proof that the match is still valid.

18.2 CURRENT PLAYER STATUS - BOTH PLAYERS

Freshly research BOTH players for:

- current ranking
- ranking movement/trajectory
- last 5 matches
- last 10 matches when obtainable
- current winning/losing streak
- same-surface recent record
- recent tournament results
- recent quality wins
- recent bad losses
- opponent quality
- current workload
- recent sets/games/minutes when obtainable
- travel
- qualifying workload
- late finishes
- injury
- illness
- retirement
- withdrawal
- medical timeout
- return from injury
- current match fitness

Do not use generic reputation as a substitute for current status.

18.3 FAVORITE LOSS AUTOPSY - FRESH WEB RESEARCH

Page 30

Research the favorite's recent losses.

For each relevant recent loss, determine when supported:

- opponent
- opponent quality
- surface
- tournament level
- score
- whether the favorite won or lost Set 1
- whether the favorite held a lead
- whether the favorite collapsed after leading
- opponent style
- favorite serve performance

- favorite return performance
- first-serve/second-serve issues
- break-point performance
- break consolidation problems
- tiebreak problems
- short-rally problems
- long-rally problems
- fatigue/injury/physical issue
- tactical pattern responsible for the loss

The purpose is to extract REPEATABLE FAVORITE FAILURE MODES.

Do not merely state:

"Favorite recently lost."

Determine HOW and WHY the favorite lost.

If detailed match statistics are unavailable, state which loss-autopsy fields are unavailable.

18.4 UNDERDOG WIN AUTOPSY - FRESH WEB RESEARCH

Research the predicted loser's recent wins.

For each relevant recent win, determine when supported:

- opponent
- opponent quality
- surface
- tournament level
- score
- whether opponent was favored/higher ranked
- whether the underdog lost Set 1 and recovered
- whether the underdog protected a lead

Page 31

- serve dominance
- return pressure
- first-strike ability
- second-serve attacking
- break-point performance
- comeback ability
- tiebreak performance
- rally tolerance
- pace tolerance
- tactical pattern
- physical endurance
- closing ability

The purpose is to extract REPEATABLE UNDERDOG WINNING WEAPONS.

Do not merely state:

"Underdog has upset wins."

Determine HOW those wins were created.

18.5 FAVORITE FAILURE MODE VS UNDERDOG WEAPON COMPATIBILITY

This comparison is mandatory.

Compare:

Favorite's demonstrated failure modes

VERSUS

Underdog's demonstrated winning weapons

A favorite weakness counts as meaningful trap evidence ONLY when this specific underdog has demonstrated an ability to exploit it.

Example:

Favorite struggles against:

- big first serves
- aggressive second-serve returns
- low-rally-length tennis

Underdog demonstrates:

- strong recent first-serve effectiveness
- aggressive return positioning
- strong 0-4 shot performance

Result:

MATCHUP-SPECIFIC TRAP PATHWAY

Generic favorite weakness without opponent compatibility receives reduced weight.

Page 32

Generic underdog strength that does not attack a favorite weakness also receives reduced weight.

The audit must distinguish:

- theoretical pathway
- demonstrated pathway
- same-surface demonstrated pathway
- direct-H2H demonstrated pathway

18.6 SAME-SURFACE DIRECT EVIDENCE

Freshly search for:

- overall H2H
- same-surface H2H
- recent same-surface meetings
- scorelines
- set-by-set pattern
- match statistics
- winners/errors when available
- first/second serve performance
- return performance
- break-point performance
- tiebreaks
- momentum changes
- retirement/injury context
- whether the meeting occurred before a major player-level change

Same-surface recent H2H receives substantially more weight than old, different-surface H2H.

A direct H2H result is not automatically decisive. Determine whether the mechanism of that result remains relevant to the current matchup.

18.7 COMMON-OPPONENT ONLINE RESEARCH

Freshly search recent common opponents.

Compare:

- result
- scoreline
- surface
- tournament level
- recency
- opponent ranking/strength
- opponent physical condition
- whether the common opponent played a similar tactical match
- whether the comparison is actually meaningful

Page 33

Do NOT treat common-opponent results as transitive proof.

Use them as contextual evidence only.

18.8 CURRENT MARKET RESEARCH

Search CURRENT odds during the CURRENT audit run.

Attempt multiple independent sportsbooks and/or reputable odds aggregators when available.

Record when obtainable:

- source/sportsbook
- favorite price
- underdog price
- implied probability
- no-vig probability
- opening price
- current price
- line movement
- disagreement between books
- whether market direction agrees with Matrix
- whether market movement strengthens the underdog trap case

If current odds cannot be reliably obtained, explicitly record:

CURRENT MARKET: UNAVAILABLE

Do NOT substitute Matrix Market Consensus for a fresh market check.

Do NOT use stale odds as current odds.

18.9 CURRENT NEWS SEARCH

Freshly search BOTH player names individually.

Look for:

- injury
- illness
- fatigue
- comments about physical condition
- travel problems
- scheduling issues
- recent retirement
- recent withdrawal
- medical timeout
- emotional/mental comments when reliably reported
- coaching changes

- equipment changes
- unusually heavy workload
- return from injury
- comments about surface adaptation
- current tournament comments

Separate:

- verified reporting
- player/coach statements
- statistical inference
- speculation

Speculation receives no decision weight unless independently supported.

18.10 SOURCE DIVERSITY REQUIREMENT

Do not call external research complete based on one website when other reliable evidence is reasonably available.

Attempt to use multiple source types:

1. official tournament/tour/draw source
2. results/statistics source
3. reputable current-news source
4. market source when available

One source may support multiple facts, but a single source must not constitute the entire external audit when other reliable evidence is available.

Source quality must be considered.

Prefer:

- official ATP/WT/ITF/tournament sources
- reputable statistical databases
- reputable news organizations
- established odds/market sources

Low-quality aggregation or scraped pages should not override stronger primary/reputable evidence.

18.11 FRESHNESS REQUIREMENT

For CURRENT:

- injury
- illness
- draw
- matchup validity
- market
- form
- ranking

- workload
- physical condition
- tournament status

prefer information from the current tournament/week whenever available.

Older information may be used for:

- H2H
- historical surface tendencies

- long-term style

- career patterns

- historical failure modes

but it must be labeled historical.

Do not present old information as current.

18.12 RESEARCH RECENCY LABELS

Every external item that materially affects the verdict should be classified as one of:

- CURRENT TOURNAMENT

- CURRENT WEEK

- RECENT - LAST 30 DAYS

- RECENT - LAST 90 DAYS

- HISTORICAL

Current information receives more weight for physical condition, form, workload, and market.

Historical information may receive strong weight when it demonstrates stable style or same-surface matchup compatibility.

18.13 ONLINE FAVORITE FAILURE LEDGER

For every favorite, record:

- Recent losses searched: YES/NO

- Relevant losses found: [N]

- Same-surface losses found: [N]

- Failure mode 1:

- Failure mode 2:

- Failure mode 3:

- Physical failure evidence:

- Closing failure evidence:

- Serve failure evidence:

- Return failure evidence:

- Tactical failure evidence:

- Reliability of failure-mode evidence: HIGH/MEDIUM/LOW

Page 36

- Current relevance: HIGH/MEDIUM/LOW

If no repeatable failure mode is found, explicitly record:

NO STRONG REPEATABLE FAVORITE FAILURE MODE FOUND

18.14 ONLINE UNDERDOG WEAPON LEDGER

For every predicted loser, record:

- Recent wins searched: YES/NO

- Relevant wins found: [N]

- Same-surface wins found: [N]

- Higher-ranked/favored opponents beaten:

- Weapon 1:

- Weapon 2:

- Weapon 3:

- Serve-based pathway:

- Return-based pathway:

- Rally-based pathway:

- pressure/closing pathway:

- fatigue/endurance pathway:
- tactical pathway:
- Reliability of weapon evidence: HIGH/MEDIUM/LOW
- Current relevance: HIGH/MEDIUM/LOW

If no meaningful underdog weapon is found, explicitly record:

NO STRONG DEMONSTRATED UNDERDOG WIN PATHWAY FOUND

18.15 FRESH ONLINE TRAP PATHWAY COUNT

After fresh research, count ONLY independent, documented pathways.

Do not double-count correlated evidence.

Example:

- strong first serve
- high ace count
- high first-serve points won

may represent ONE serve-dominance pathway, not three independent pathways.

Classify each pathway:

1. PROVEN DIRECT PATHWAY

- demonstrated against this favorite, preferably same surface

2. PROVEN COMPARABLE PATHWAY

Page 37

- demonstrated against similar/high-quality opponents

3. THEORETICAL PATHWAY

- stylistically plausible but not demonstrated strongly enough

Only proven direct/comparable pathways receive full trap weight.

18.16 FRESH ONLINE STRESS TEST

After online research, re-run the favorite verdict under adverse but plausible assumptions.

At minimum test:

Stress Test A - Remove strongest favorite model family

Would the favorite still be favored?

Stress Test B - Give greater weight to current form

Would the predicted loser become competitive?

Stress Test C - Give greater weight to same-surface evidence

Would the predicted loser become competitive?

Stress Test D - Give greater weight to current physical/news evidence

Would the favorite remain stable?

Stress Test E - Give greater weight to current market

Would the winner change?

Stress Test F - Assume one-break/tiebreak match

Does the underdog's pathway become substantially stronger?

Record:

- winner under each stress test
- whether verdict flips
- how many plausible stress tests flip the favorite

A favorite that flips under several plausible stress tests receives a major trap-risk increase.

18.17 CURRENT MARKET VS MATRIX MARKET SEPARATION

The audit must explicitly separate:

MATRIX MARKET CONSENSUS

from

FRESH CURRENT ONLINE MARKET

These are not the same evidence source.

Record both separately.

If they agree:

MANDATORY MARKET CONSENSUS ENFORCEMENT -

DISAGREEMENT / TRAP AUDIT

This page is part of the audit, not an optional appendix. It strengthens Sections 18.8 and 18.17 and controls every market-read, trap-market, Market-Veto, and final KEEP/DOWNGRADE/REMOVE/FLIP/PASS decision.

1. Sportsbook source minimum. For every current match, attempt at least 3 independent sportsbooks when available. Prefer one sharp/reference or high-liquidity book plus two mainstream regulated books. An odds aggregator may help locate prices, but it does not count as multiple independent books unless the underlying books are individually identified.
2. Same-market / same-time rule. Capture both players from the same match-winner market and the same book at approximately the same timestamp. Record sportsbook, timestamp, favorite price, underdog price, and source URL. Suspended, stale, one-sided, or mismatched-round prices are not valid consensus inputs.
3. De-vig each sportsbook separately. Convert both sides to raw implied probability, then compute the book-level no-vig favorite probability as $qFav / (qFav + qDog)$. Build the live sportsbook consensus from the median no-vig probability across valid books; also record low/high range and dispersion.
4. Prediction markets are separate. Kalshi and other prediction markets must be recorded as a separate evidence family. Never treat Kalshi American-style display odds as ordinary sportsbook moneylines, never blend them into sportsbook consensus, and never use them as a substitute for a two-sided sportsbook quote.
5. Kalshi coherence check. When Kalshi is used, prefer the underlying contract price/probability and liquidity/spread information. If its displayed American-style numbers produce an extreme synthetic overround, cannot be reconciled with complementary contract probabilities, or materially conflict with the underlying contract price, label the display NON-COMPARABLE / DISPLAY-STRUCTURE WARNING and give it no sportsbook-consensus weight.
6. Market integrity thresholds. A sportsbook-consensus vs Matrix-market gap of 5+ percentage points requires investigation; 8+ pp is a MATERIAL MARKET CONFLICT; 12+ pp is a CRITICAL MARKET CONFLICT and blocks Green/Keep until explained. Apply the same 5/8/12 pp ladder to Kalshi vs sportsbook consensus, but treat the result as a prediction-market discrepancy rather than sportsbook evidence.
7. Book-dispersion warning. If valid sportsbook no-vig probabilities span 8+ pp, mark MARKET DISPERSION HIGH and do not claim a precise consensus. Investigate stale lines, limits/liquidity, player identity, tournament/round, withdrawals, and news before assigning decision weight.
8. Source-count downgrade. Three or more valid books = COMPLETE. Two valid books = PARTIAL and reduced market weight. One valid book = SINGLE-SOURCE ONLY and cannot independently unlock Green. Zero reliable books = CURRENT SPORTSBOOK MARKET UNAVAILABLE.
9. No market double-counting. Matrix Market Consensus, live sportsbook consensus, and prediction-market information must remain three distinct fields. Agreement may strengthen confidence; disagreement must be explained. Do not count multiple books as multiple independent tennis signals; they form one Market evidence family.
10. Market cannot be silently ignored. A material/critical conflict must appear in the contradiction ledger, trap score, stress tests, removal case, and final decision rationale. If unresolved, downgrade/remove/pass rather than certifying Green.

Mandatory per-match market ledger

Field

Required entry

Sportsbooks

Book 1 / Book 2 / Book 3+; prices; timestamp; URL

No-vig consensus

Median; low-high range; dispersion; valid source count

Matrix market

Separate Matrix Market Consensus probability

Prediction market

Kalshi/other contract probability, liquidity/spread; separate

Conflict flags

Matrix-vs-sportsbook pp gap; prediction-vs-sportsbook pp gap

Integrity status

COMPLETE / PARTIAL / SINGLE-SOURCE / UNAVAILABLE

Decision effect

CONFIRMS / NEUTRAL / WEAKENS / MATERIAL CONFLICT / CRITICAL VETO

DISAGREEMENT/TRAP OVERRIDE: A CRITICAL market conflict or non-comparable Kalshi display cannot be used to support the favorite. The audit must investigate the underdog pathway and present the unresolved conflict before final classification.

Page 38

- agreement may strengthen the favorite/underdog case.

If they disagree:

- investigate timing, line movement, injury/news, and possible stale Matrix inputs.

Never report Matrix market data as if it were a fresh live market check.

18.18 FRESH ONLINE RESEARCH LEDGER - REQUIRED FOR EVERY MATCH

For EVERY matchup report:

- Favorite fresh research: COMPLETE / INCOMPLETE

- Underdog fresh research: COMPLETE / INCOMPLETE

- Current draw verified: YES / NO

- Current tournament verified: YES / NO

- Current surface verified: YES / NO

- Current round verified: YES / NO / UNAVAILABLE

- Recent favorite losses researched: YES / NO

- Favorite failure modes extracted: YES / NO / NONE FOUND

- Recent underdog wins researched: YES / NO

- Underdog winning weapons extracted: YES / NO / NONE FOUND

- Same-surface evidence researched: YES / NO

- H2H researched: YES / NO

- Common opponents researched: YES / NO

- Current injury/status researched - favorite: YES / NO

- Current injury/status researched - underdog: YES / NO

- Current workload researched - favorite: YES / NO

- Current workload researched - underdog: YES / NO

- Current market researched: YES / UNAVAILABLE

- Matrix market separated from live market: YES / NO

- Current news researched: YES / NO

- Multiple source types attempted: YES / NO

- Source freshness classified: YES / NO

- Failure-mode compatibility test completed: YES / NO

- Fresh online stress test completed: YES / NO

- Independent trap pathways counted: YES / NO

18.19 HARD STOP - FRESH ONLINE RESEARCH

Previous research DOES NOT count toward the Part 18 ledger.

Every required field must be newly researched during the CURRENT audit execution.

If mandatory fresh research has not been performed for BOTH players, the matchup cannot receive a final:

- KEEP

Page 39

- DOWNGRADE

- REMOVE

- FLIP

- confirmed trap

- safest-favorite certification

It must instead say:

INCOMPLETE FRESH-ONLINE AUDIT - missing: [specific online research item/source].

If current market is genuinely unavailable after a documented search attempt, it may be marked UNAVAILABLE rather than NO. That alone does not automatically block completion, provided the audit clearly states the limitation and does not invent or substitute odds.

If current draw/match validity cannot be verified, the matchup CANNOT receive final certification.

18.20 NO PREVIOUS-RESEARCH SUBSTITUTION

The following DO NOT satisfy Part 18:

- earlier assistant web searches

- previous audit external-validation text

- prior Green verification

- cached injury notes

- remembered rankings

- old odds

- old draw information

- previous H2H summaries

- previous news summaries

Those materials may be used as leads for what to re-check, but the relevant facts must be freshly verified during the current run.

18.21 BATCH CERTIFICATION ADDITION

The final AUDIT CERTIFICATION must now also include:

- Fresh favorite web audits completed: [N/N]

- Fresh underdog web audits completed: [N/N]

- Current draw validations completed: [N/N]

- Favorite loss autopsies completed: [N/N]

- Underdog win autopsies completed: [N/N]

- Same-surface online checks completed: [N/N]

- H2H online checks completed: [N/N]

- Common-opponent online checks completed: [N/N]

- Injury/status online checks completed: [N/N]

- Current market searches attempted: [N/N]

- Fresh news searches completed: [N/N]

- Failure-mode compatibility tests completed: [N/N]

- Fresh online stress tests completed: [N/N]

- Source-diversity requirement satisfied: [N/N or documented exceptions]

Only when the original completion ledger AND this Fresh Online Research ledger satisfy all mandatory requirements may the assistant write:

FULL DISAGREEMENT/TRAP AUDIT COMPLETE - NO REQUIRED STEP SKIPPED.

18.22 ENFORCEMENT PRECEDENCE

Part 18 is mandatory and works together with Part 16.

If any earlier section can be interpreted as allowing previous/cached external research to satisfy current online validation, Part 18 overrides that interpretation.

A disagreement/trap audit is not fully executed merely because online research exists somewhere in a previous report.

The research must be FRESH, MATCH-SPECIFIC, PLAYER-SPECIFIC, OPPONENT-SPECIFIC, and performed during the CURRENT audit run.

FINAL ENFORCEMENT PRINCIPLE

The disagreement audit must answer two separate questions with current evidence:

1. WHY CAN THE FAVORITE LOSE?

2. DOES THIS SPECIFIC PREDICTED LOSER CURRENTLY POSSESS THE TOOLS TO MAKE THAT LOSS HAPPEN?

The first question requires a fresh favorite loss autopsy.

The second requires a fresh underdog win autopsy.

A trap verdict requires BOTH.

If either side has not been freshly researched, the trap audit is incomplete.

ADDED DISAGREEMENT RULES - FALSE-FAVORITE ESCALATION

These rules are mandatory additions to the Disagreement/Trap Audit. Their purpose is to force removal or downgrade review when a historically strong favorite faces a rapidly improving, match-sharp, or under-ranked opponent.

FALSE-FAVORITE ESCALATION GATE - GREEN LOCK

If the favorite has meaningful inactivity or no current-condition evidence AND the opponent shows rapid improvement, ranking lag, recent elite wins, or unusually strong opponent-adjusted form, GREEN and DOUBLE GREEN are LOCKED until the escalation below is completed.

1. Current-Strength Revaluation

Estimate both players' present strength separately from career/name value using recent opponent-adjusted results, surface performance, serve/return trajectory, ranking/Elo velocity, and quality of opposition.

2. Rapid-Riser / Ranking-Lag Override

If the underdog's recent level is materially ahead of ranking/Elo, long-term ratings cannot dismiss the underdog. Quantify the gap between trailing baseline and recent current-strength estimate.

3. Recent Elite-Win Escalation

Investigate recent high-quality wins for repeatability using underlying points/games, serve/return, opponent quality, surface, and follow-up performance.

4. Favorite Rust / Layoff Penalty

Meaningful inactivity widens the favorite's floor and blocks clean confidence based primarily on historical strength. Compare layoff length, match sharpness, current-surface adaptation, and first-match-back history when reconstructable.

5. No Current-Tournament Evidence Penalty

If the underdog has current-condition evidence and the favorite has none, increase favorite uncertainty and require affirmative current-strength support before Green.

6. Rust x Momentum Cross-Feature

Directly test returning established favorite versus match-sharp / rapidly improving opponent. This interaction can create trap risk even when each factor alone is moderate.

7. Underdog Ceiling Revaluation

Recent elite-quality performances update the underdog's realistic upper range. Recalculate favorite-floor vs underdog-ceiling overlap.

8. Favorite Floor Revaluation

Layoff, volatility, uncertain current condition, or recent decline widens the favorite's lower tail. Recalculate distribution overlap.

9. Information-Asymmetry Penalty

If present-level evidence is much richer for the underdog than the favorite, do not let historical certainty substitute for current evidence.

10. Legacy / Name Bias Test

Ignore names, titles, reputation, and career accomplishments. Re-evaluate using recent 10-20 matches, surface, opponent quality, serve/return trajectory, workload, and current-condition evidence. If the color changes, flag LEGACY/NAME BIAS RISK.

11. Market-Veto Investigation

Large unexplained market disagreement blocks Green. Investigate line level/movement, source quality, current availability news, and whether the market is pricing a present-level risk the historical metrics miss.

12. Probability Compression

When present-level uncertainty is high, reject extreme certainty. Widen the probability interval and test whether the pick remains favorable under conservative current-strength assumptions.

13. Mandatory Removal Question

Before any Green list is finalized, answer: What is the strongest evidence that this favorite should be REMOVED? If the audit cannot convincingly defeat that case with independent reliable evidence, remove or downgrade the favorite.

14. Adversarial Breakthrough Case

Construct the strongest evidence-based case that the improving underdog is already better than the long-term rating suggests. Test it against surface, opponent adjustment, serve/return, current tournament, market, and sample reliability.

15. Green Burden-of-Proof Rule

In an escalation match, absence of proof for the underdog is not proof for the favorite. The favorite must affirmatively demonstrate current superiority after all adjustments.

16. False-Green Historical Feedback

Every Green/Double-Green loss becomes a structural failure case. Store the pre-match warning configuration and use it as a future downgrade trigger when the same pattern recurs.

17. No Future Leakage

Only pre-match information may explain whether the pick should have been removed. In-match illness, vomiting, injury, medical timeout, or performance cannot be used retroactively unless publicly known before first serve.

18. Escalation Exit Standard

GREEN unlocks only if the favorite survives current-strength revaluation, trajectory differential, floor/ceiling overlap, market investigation, adversarial underdog case, and removal-question review. Otherwise downgrade to YELLOW/RED/PASS as supported.

DANGEROUS UNDERDOG SCREEN

MANDATORY FALSE-FAVORITE ESCALATION - RUN BEFORE VALIDATING THE FAVORITE

Purpose: detect an under-ranked, rapidly improving, match-sharp, or matchup-dangerous opponent before historical reputation, ranking, Elo, or Matrix probability can push the favorite onto the Green list.

Five-Dimension Screen

CURRENT LEVEL - Is the underdog materially better now than long-term ranking/Elo indicates?

TRAJECTORY - Is the underdog improving quickly across ranking, opponent-adjusted performance, serve, return, and quality of opposition?

PROVEN CEILING - Has the underdog already produced matchup-winning performance against strong opposition, and is that ceiling repeatable?

MATCHUP ACCESS - Does the underdog possess one or more credible, opponent-specific pathways that attack the favorite's documented weaknesses?

FAVORITE EXPOSURE - Is the favorite's historical edge less trustworthy today because of inactivity, decline, surface, current-condition uncertainty, weak underlying performance, or matchup vulnerability?

3-of-5 GREEN LOCK

If 3 or more dimensions trigger STRONGLY, GREEN and DOUBLE GREEN are unavailable until the favorite completes the full

Dangerous Underdog Escalation. Absence of proof for the underdog is not proof for the favorite.

Required Screen Output

- Current Level: LOW / MODERATE / HIGH threat
- Trajectory: LOW / MODERATE / HIGH threat
- Proven Ceiling: LOW / MODERATE / HIGH threat
- Matchup Access: LOW / MODERATE / HIGH threat
- Favorite Exposure: LOW / MODERATE / HIGH threat
- Strong triggers: [0-5]
- GREEN LOCK: YES / NO
- Strongest evidence for removal: [state it before defending the favorite]

ADDITIONAL DISAGREEMENT ENFORCEMENT - RULES

/ SIGNALS 1-70

All items use public observations or transparently reconstructed values. DIRECT, RECONSTRUCTED, PARTIAL, UNAVAILABLE, and EXCLUDED must remain distinguished. Correlated metrics are not independent votes.

1. Underdog Threat Score - mandatory - Build a separate threat score without using the Matrix probability: recent opponent-adjusted performance, ranking velocity, surface performance, elite wins, serve/return improvement, current-tournament performance, quality of losses, volatility/upside, common-opponent performance, and market strength.
2. Favorite Vulnerability Score - mandatory - Separately measure inactivity, recent bad losses, declining serve/return, ranking-performance disconnect, closing failures, break-lead surrender, tiebreak dependence, weak-win frequency, workload, and surface/current-condition uncertainty.
3. Recent-Level Elo - Reconstruct strength over multiple windows such as 52, 26, 12, and approximately 6 weeks where inputs permit. Compare with long-term Elo and label all reconstructed ratings.
4. Ranking Velocity + Acceleration - Measure ranking change and whether the rate of improvement/decline itself is accelerating.
5. Opponent-Quality Acceleration - Test whether recent performance is occurring against progressively stronger opposition.
6. Upset Capability Index - Measure recent performance against Top 10/20/50/100, higher-rated opponents, and betting favorites where historical pricing is available.
7. Upset Recency Index - Weight recent high-quality upsets more heavily than old upsets while preserving sample reliability.
8. Near-Upset Index - Capture competitive losses to elite opponents using sets, tiebreaks, game/point dominance, break differential, and closeness.
9. Hidden Improvement Index - Detect rising point%, hold%, break%, second-serve, return, and opponent-adjusted performance even when W/L remains mediocre.
10. Quality-of-Loss Trend - Determine whether recent losses are increasingly restricted to stronger opponents and whether underlying performance in those losses is improving.
11. Favorite Bad-Loss Frequency - Count losses to materially lower-ranked/lower-rated opponents over multiple recency windows and on the relevant surface.
12. Favorite Escape Frequency - Count wins where the favorite survived despite weak underlying point/game/break performance or repeated pressure exposure.
13. Expected-Wins vs Actual-Wins Gap - Compare recent actual W/L with expected results from transparent pre-match strength or point/game performance estimates.
14. Opponent-Adjusted Point Dominance Trend - Track point-level dominance after opponent-strength adjustment and compare trend direction for both players.
15. Serve Residual Trend - Actual serve performance minus expected serve performance against the returners faced.
16. Return Residual Trend - Actual return performance minus expected return performance against the servers faced.
17. Second-Serve Survival Trend - Track second-serve points won and exposure after opponent adjustment, especially against aggressive return profiles.

18. Break Creation Trend - Track break-point opportunities generated per return game and its recent/surface/opponent-adjusted trajectory.
19. Break Prevention Trend - Track break-point opportunities conceded per service game and trajectory.
20. Pressure Opportunity Differential - Compare how frequently each player creates and faces high-leverage break situations, not merely conversion percentages.
21. Match-Control Without Scoreline - Reconstruct control from points, service points, return points, breaks, and break opportunities while deliberately ignoring the final result.
22. Performance Stability vs Strong Opposition - Measure whether performance deteriorates sharply or remains stable as opponent quality increases.
23. Opponent Strength Elasticity - Estimate how much point/serve/return performance changes as opponent strength rises.
24. Surface Strength Elasticity - Estimate whether the matchup gap expands or compresses on today's surface relative to overall performance.
25. Favorite Advantage Compression - Measure how much the favorite's edge shrinks through overall -> surface -> recent -> opponent-adjusted -> current-tournament filters.
26. Underdog Threat Expansion - Measure whether the underdog's estimated strength rises as stale/irrelevant historical information is removed.
27. Peak-Recent-Performance Frequency - Count how often the underdog has actually reached an elite or matchup-winning performance band recently.
28. Ceiling Repeatability - Distinguish one outlier peak from repeated high-level performances over the recent sample.
29. Performance Change-Point Detection - Use transparent statistical change-point or segmented-trend methods where sample permits to identify meaningful shifts in current level.
30. EWMA Current Form - Exponentially weight recent point/game/serve/return performance so the newest evidence matters more without discarding older observations.
31. Separate Serve EWMA and Return EWMA - Maintain independent exponentially weighted serve and return trajectories.
32. Current-Level Confidence Interval - Every reconstructed current-strength estimate must include uncertainty/range and sample basis.
33. Probability Range Instead of Single Probability - Report a plausible reconstructed probability interval, not only a single point estimate.
34. Worst-Reasonable-Case Simulation - Stress-test lower-bound favorite assumptions against upper-bound underdog assumptions. If modest changes flip the winner, Green is fragile.
35. Monte Carlo Parameter Uncertainty - Sample plausible serve/return/hold parameters from uncertainty ranges instead of fixing one estimate.
36. Upset Monte Carlo - Estimate how often plausible favorite-floor/underdog-ceiling combinations produce the upset.
37. Path-to-Upset Simulation - Estimate plausible upset routes: straight sets, decider, break-heavy, tiebreak, second-serve attack, closing failure, etc.
38. Multiple Independent Upset Routes - Count independent credible pathways by which the underdog can win. Multiple routes materially increase trap risk.
39. Network Strength Rating - Build a transparent player-opponent graph from recent public results to identify players whose ranking may lag their competitive network position.
40. Temporal Network Rating - Recency-weight the network so current opponent relationships matter more than stale ones.
41. Second-Degree Opponent Strength - Evaluate not only who a player beat, but the recent strength/results of those opponents' opponents.
42. Quality-Win Network Density - Test whether strong results form a connected pattern against genuinely strong competition rather than isolated anomalies.
43. Common-Opponent Recency - Weight recent same-surface common opponents more heavily than old comparisons.
44. Common-Opponent Performance Residual - Compare each player's performance against shared opponents relative to expected performance.
45. Archetype-Neighborhood Analysis - Find statistically similar recent opponents using observable surface, strength, serve/return, rally, and style variables rather than broad labels alone.
46. Nearest-Opponent Match Search - Construct a transparent similarity vector and retrieve the closest historical opponent profiles for each player.
47. Favorite Failure Twins - Find past favorite losses/struggles against opponents most similar to today's underdog.
48. Underdog Success Twins - Find underdog wins/near-wins against opponents most similar to today's favorite.
49. Twin Outcome Rate - Calculate outcomes across the comparable-match sample instead of relying on anecdotal similarities.
50. Twin Reliability Gate - No historical-twin conclusion without sample size, similarity strength, recency, and surface relevance.
51. Superstar / Name-Value Guard - Automatically escalate historically elite favorites when the opponent's current level is rising materially faster than

ranking recognition.

52. Career-vs-Current Strength Divergence - Quantify the gap between long-term/career strength and current reconstructed strength for both players.

53. Historical Reputation Contribution - Recalculate the favorite's advantage after excluding evidence older than 12 months where possible; report how much the edge changes.

54. Favorite Recognition Bias Alarm - If career evidence strongly favors the favorite but recent/surface/opponent-adjusted evidence is close, force downgrade review.

55. Freshness Asymmetry Alarm - If most favorite-supporting evidence is old while underdog-supporting evidence is recent, lock Green pending escalation.

56. Current Evidence Coverage Ratio - Report what percentage of decision evidence comes from last 4/8/12/26 weeks.

57. Independent Evidence Effective Count - Reduce correlated statistics into effective independent evidence families rather than counting every metric as a separate vote.

58. Contradiction Concentration - Distinguish many contradictions from one correlated family versus contradictions across multiple independent families.

59. Underdog Evidence Diversity Score - Count and weight independent evidence families supporting the upset case.

60. Favorite Evidence Diversity Score - Require comparable diversity in the favorite's case before Green.

61. Break-After-Break Momentum Profile - From game/point logs, reconstruct break consolidation and immediate rebreak patterns for both players.

62. Interruption Sensitivity Proxy - Where logs and public interruption data support it, compare observable performance before/after delays or interruptions; do not infer psychology.

63. Error Pressure Profile - When charting exists, distinguish winners, forced errors, unforced errors, and double faults under opponent pressure.

64. Error Spike Vulnerability - Identify whether favorite losses repeatedly contain abrupt double-fault/unforced-error deterioration under return pressure.

65. Underdog Error-Induction Ability - Determine whether the underdog repeatedly causes stronger opponents to produce worse serve/shot outcomes than their baselines.

66. Match-State Resilience - Reconstruct performance after being broken, losing a set, failing to serve out a set, or losing a tiebreak.

67. Favorite Shock Response - Specifically test favorite performance after losing the first break/set against lower-rated opposition.

68. Underdog Lead Protection vs Better Players - Measure whether the underdog can consolidate and close leads against higher-rated opponents.

69. Upset Conversion Rate - Among matches where the underdog reaches a break/set lead against materially stronger players, measure how often the upset is completed.

70. Dangerous-Underdog Gate - If rapid-riser/ranking-lag + current-strength improvement + credible elite-opponent performance trigger together, Green is unavailable until the Dangerous Underdog Screen is completed.

MANDATORY EXIT TEST

The favorite cannot unlock GREEN merely because it remains the higher-ranked or historically stronger player. GREEN unlocks only after the audit defeats the strongest removal case using multiple independent, reliable, current, surface-relevant, opponent-adjusted evidence families and the conclusion survives uncertainty/sensitivity testing.

DATA-COVERAGE RULE

When advanced point/shot/charting data are unavailable for a player, mark the affected metric PARTIAL or UNAVAILABLE and continue. Missing charting data must increase uncertainty; it must never be replaced with invented observations or silently interpreted as evidence for the favorite.

INTEGRATED VETO & HANDOFF ARCHITECTURE

Mandatory addition to the Disagreement/Trap Audit. This section gives the disagreement system explicit veto/handoff authority inside the combined workflow.

1. Final Combined Decision Gate

A candidate reaches the Final Green List only if the Verification Audit approves it AND the Disagreement/Trap Audit fails to falsify it AND the Dangerous Underdog Screen is clear AND the final stress/removal tests are passed. Failure at any required stage blocks Green.

2. Disagreement Veto Power

A RED/PASS, unresolved CRITICAL contradiction, or active Dangerous Underdog GREEN LOCK from the Disagreement Audit automatically prevents Green regardless of the Matrix probability or Verification approval.

3. Hard Downgrade Triggers

Specified combinations must mechanically downgrade rather than merely warn. Example: rapid-riser/ranking-lag opponent + credible recent elite performance + favorite inactivity/current-level uncertainty locks Green until escalation is passed.

4. Minimum Independent-Evidence Requirement

Do not approve from raw metric count. Green requires at least 4 genuinely independent reliable evidence families supporting the candidate after contradictions; Double Green requires at least 6. Correlated metrics collapse into one effective family.

5. Contradiction Budget

Double Green: zero unresolved CRITICAL contradictions and no more than one unresolved MAJOR contradiction. Green: zero unresolved CRITICAL contradictions and only limited MAJOR contradictions that are explicitly defeated. Otherwise downgrade.

6. Evidence Freshness Score

Report how much decisive evidence comes from the last 4/8/12/26 weeks and how much is older. Stale evidence cannot carry the same confidence as fresh, surface-relevant, opponent-adjusted evidence.

7. Historical Strength vs Current Strength Split

Finish every match with separate Historical Strength and Current Estimated Strength assessments. A material divergence is a mandatory warning and triggers deeper current-level review.

8. Probability Stress Range

Report Base, Conservative, and Optimistic reconstructed probability/range. If reasonable conservative assumptions push the favorite near 50% or flip the winner, Green is fragile and must be downgraded unless independently resolved.

9. Removal-First Final Stage

Before approval, construct the strongest evidence-based REMOVE CASE against the Matrix winner. The candidate must defeat that case using independent reliable evidence. Failure to defeat it blocks Green.

10. False-Green Learning Ledger

Every Green/Double-Green loss receives a pre-match failure classification using only information available before first serve: rapid riser missed, stale rating, weak-win inflation, current decline, matchup collision, market contradiction, inactivity, etc. Recurrent structures become future downgrade triggers.

11. Calibration by Audit Color

Track actual outcomes separately for Double Green, Green, Yellow, Red, and PASS/removed candidates. If higher colors do not materially outperform lower colors over adequate samples, tighten thresholds.

12. Audit Confidence Is Not Win Probability

Always report Audit Confidence separately from estimated win probability. A high estimated probability with incomplete/low-quality evidence cannot receive equally high audit confidence.

13. Minimum Coverage Threshold

Report Audit Coverage %. Double Green requires very high coverage of CRITICAL evidence families; Green requires sufficient critical coverage. Missing supplemental metrics alone do not block approval.

14. Critical-Metric Coverage

Classify audit fields as CRITICAL / IMPORTANT / SUPPLEMENTAL. Missing current surface form, serve/return, opponent quality, availability, or present-strength evidence matters far more than missing an obscure tactical statistic.

15. Effective Independent Evidence Count

Final output must report the number of effective independent evidence families after correlated-data compression. Do not allow dozens of correlated metrics to create false unanimity.

16. Mandatory Strongest-Underdog Table

For every proposed Green, list the underdog's 5 strongest pre-match arguments, reliability/sample, evidence family, and whether each was DEFEATED / UNRESOLVED / SUPPORTS REMOVAL.

17. Final Rejection Reason

Every removed/downgraded favorite receives one concise human-readable reason identifying the decisive structural risk rather than merely listing

metrics.

Disagreement-Specific Handoff Rule

The Disagreement Audit must explicitly return one of: CLEAR / CAUTION / GREEN LOCK / RED VETO / PASS VETO. CLEAR permits progression; CAUTION requires final stress review; GREEN LOCK/RED VETO/PASS VETO blocks Green.

Dangerous Underdog Handoff Rule

A 3-of-5 strong Dangerous Underdog trigger cannot be ignored by the combined decision gate. The favorite must pass the full escalation exit standard before the lock can be cleared.

FINAL INTEGRATED WORKFLOW

MANDATORY WORKFLOW: Tennis Matrix Candidate -> Verification Audit -> Disagreement/Trap Audit -> Dangerous Underdog Screen -> Stress Test -> Removal Case -> FINAL GREEN LIST

No candidate may bypass a stage. A blocked stage must be resolved or the candidate remains downgraded/removed. The purpose is filtering, not forcing a winner.